

YOU ARE HERE: [HOME](#) > [JDK VERSIONS](#) > [JAVA 7 FEATURES](#) > [FORK/JOIN FRAMEWORK TUTORIAL: FORKJOINPOOL EXAMPLE](#)

Fork/Join Framework Tutorial: ForkJoinPool Example

🕒 MAY 27, 2014 👤 LOKESH 💬 4 COMMENTS

[Follow](#)

2.6k

+33

[Recommend this on Google](#)

The effective use of parallel cores in a Java program has always been a challenge. There were few home-grown frameworks that would distribute the work across multiple cores and then join them to return the result set. Java 7 has incorporated this feature as a **Fork and Join framework**.

Selenium Visual Testing

Add Automated Visual Testing to your Selenium tests in minutes



Basically the **Fork-Join breaks the task at hand into mini-tasks** until the mini-task is simple enough that it can be solved without further breakups. It's like a **divide-and-conquer algorithm**. One important concept to note in this framework is that **ideally no worker thread is**

CA LISA® Development

Shorten Application Release Times. Download White Paper Here for More!



CONNECT WITH ME

**Lokesh Gupta**[Follow](#)


2 699 followers

**How to do in java**

Like

10,193 people like [How to do in java](#).

idle. They implement a **work-stealing algorithm** in that idle workers **steal** the work from those workers who are busy.



Fork Join Framework

It's based on the work of Doug Lea, a thought leader on Java concurrency. Fork/Join deals with the threading hassles; you just indicate to the framework which portions of the work can be broken apart and handled recursively. It employs pseudocode (as taken from Doug Lea's paper on the subject):

```
Result solve(Problem problem) {  
    if (problem is small)  
        directly solve problem  
    else {  
        split problem into independent parts  
        fork new subtasks to solve each part  
        join all subtasks  
        compose result from subresults  
    }  
}
```

Discussion Points



Scale agile
faster,
easier,
smarter

See a Demo

YOU MAY ALSO LIKE

- 1) Core Classes used in Fork/Join Framework
 - i) `ForkJoinPool`
 - ii) `ForkJoinTask`
- 2) Example Implementations of Fork/Join Pool Framework
 - i) Implementation Sourcecode
 - ii) How it works?
- 3) Difference between Fork/Join Framework And `ExecutorService`
- 4) Existing Implementations in JDK
- 5) Conclusion

Core Classes used in Fork/Join Framework

The core classes supporting the Fork-Join mechanism are `ForkJoinPool` and `ForkJoinTask`.

Let's learn about their roles in detail.

ForkJoinPool

The `ForkJoinPool` is basically a specialized implementation of `ExecutorService` implementing the work-stealing algorithm we talked about above. We create an instance of `ForkJoinPool` by providing the target parallelism level i.e. the **number of processors** as shown below:

Java executor framework tutorial and best practices

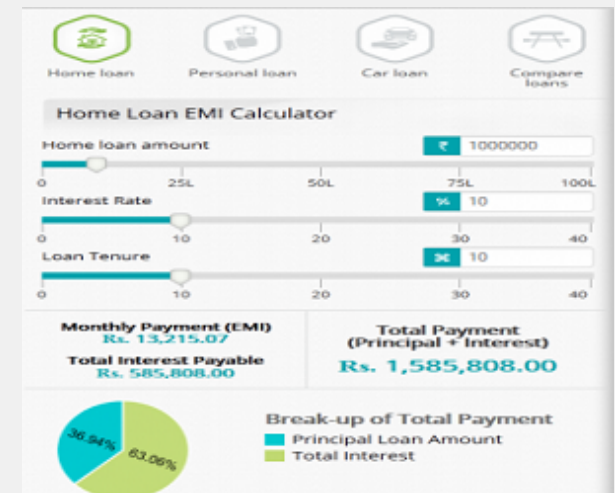
Difference between `yield()` and `join()` in threads in java?

Java 8: String join (CSV) example

Java IO : `FileFilter` Example Tutorial

Java 7 Changes, Features and Enhancements

When to use `CountDownLatch` : Java concurrency example tutorial



```
ForkJoinPool pool = new ForkJoinPool(numberOfProcessors); Where  
numberOfProcessors = Runtime.getRuntime().availableProcessors();
```



If you use a no-argument constructor, by default, it creates a pool of size that equals the number of available processors obtained using above technique.

Although you specify any initial pool size, the **pool adjusts its size dynamically in an attempt to maintain enough active threads** at any given point in time. Another important difference compared to other `ExecutorService's` is that this pool need not be explicitly shutdown upon program exit because all its threads are in daemon mode.

Amazon Kinesis
Streaming analytics
made easy.



Get Started

There are three different ways of submitting a task to the `ForkJoinPool`.

- 1) **execute() method** //Desired asynchronous execution; call its fork method to split the work between multiple threads.
- 2) **invoke() method**: //Await to obtain the result; call the invoke method on the pool.
- 3) **submit() method**: //Returns a Future object that you can use for checking status and obtaining the result on its completion.

ForkJoinTask

This is an abstract class for creating tasks that run within a `ForkJoinPool`. The `RecursiveAction` and `RecursiveTask` are the only two direct, known subclasses of `ForkJoinTask`. The only difference between these two classes is that the `RecursiveAction` does not return a value while `RecursiveTask` does have a return value and returns an object of specified type.

In both cases, you would need to implement the compute method in your subclass that performs the main computation desired by the task.

The `ForkJoinTask` class provides several methods for checking the execution status of a task. The `isDone()` method returns true if a task completes in any way. The `isCompletedNormally()` method returns true if a task completes without cancellation or encountering an exception, and `isCancelled()` returns true if the task was cancelled. Lastly, `isCompletedabnormally()` returns true if the task was either cancelled or encountered an exception.

Example Implementations of Fork/Join Pool Framework

In this example, you will learn how to use the asynchronous methods provided by the

**Amazon
Kinesis**
Process
real-time data
at massive
scale.



Try Now



`ForkJoinPool` and `ForkJoinTask` classes for the management of tasks. You are going to implement a **program that will search for files with a determined extension inside a folder and its subfolders**. The `ForkJoinTask` class you're going to implement will process the content of a folder. For each subfolder inside that folder, it will send a new task to the `ForkJoinPool` class in an asynchronous way. For each file inside that folder, the task will check the extension of the file and add it to the result list if it proceeds.

The solution to above problem is implemented in `FolderProcessor` class, which is given below:

Implementation Sourcecode

FolderProcessor.java

```
1 package forkJoinDemoAsyncExample;
2
3 import java.io.File;
4 import java.util.ArrayList;
5 import java.util.List;
6 import java.util.concurrent.RecursiveTask;
7
8 public class FolderProcessor extends RecursiveTask<List<String>> {
9     {
10         private static final long serialVersionUID = 1L;
11         //This attribute will store the full path of the folder this task is going
12         private final String path;
13         //This attribute will store the name of the extension of the files this tas
14         private final String extension;
15
16         //Implement the constructor of the class to initialize its attributes
17         public FolderProcessor(String path, String extension)
18         {
19             this.path = path;
20             this.extension = extension;
21         }
22
23         //Implement the compute() method. As you parameterized the RecursiveTask cl
24         //this method has to return an object of that type.
25         @Override
26         protected List<String> compute()
27         {
28             //List to store the names of the files stored in the folder.
```

```

29 List<String> list = new ArrayList<String>();
30 //FolderProcessor tasks to store the subtasks that are going to process
31 List<FolderProcessor> tasks = new ArrayList<FolderProcessor>();
32 //Get the content of the folder.
33 File file = new File(path);
34 File content[] = file.listFiles();
35 //For each element in the folder, if there is a subfolder, create a new
36 //and execute it asynchronously using the fork() method.
37 if (content != null)
38 {
39     for (int i = 0; i < content.length; i++)
40     {
41         if (content[i].isDirectory())
42         {
43             FolderProcessor task = new FolderProcessor(content[i].getAbsolutePath());
44             task.fork();
45             tasks.add(task);
46         }
47         //Otherwise, compare the extension of the file with the extension
48         //and, if they are equal, store the full path of the file in the list.
49         else
50         {
51             if (checkFile(content[i].getName()))
52             {
53                 list.add(content[i].getAbsolutePath());
54             }
55         }
56     }
57 }
58 //If the list of the FolderProcessor subtasks has more than 50 elements,
59 //write a message to the console to indicate this circumstance.
60 if (tasks.size() > 50)
61 {
62     System.out.printf("s: %d tasks ran.\n", file.getAbsolutePath());
63 }
64 //add to the list of files the results returned by the subtasks launched
65 addResultsFromTasks(list, tasks);
66 //Return the list of strings
67 return list;
68 }
69
70 //For each task stored in the list of tasks, call the join() method that will
71 //Add that result to the list of strings using the addAll() method.
72 private void addResultsFromTasks(List<String> list, List<FolderProcessor> tasks)
73 {
74     for (FolderProcessor item : tasks)
75     {
76         list.addAll(item.join());
77     }
78 }
79

```

```

80 //This method compares if the name of a file passed as a parameter ends with
81 private boolean checkFile(String name)
82 {
83     return name.endsWith(extension);
84 }
85 }

```

And to use above `FolderProcessor`, follow below code:

Main.java

```

1 package forkJoinDemoAsyncExample;
2
3 import java.util.List;
4 import java.util.concurrent.ForkJoinPool;
5 import java.util.concurrent.TimeUnit;
6
7 public class Main
8 {
9     public static void main(String[] args)
10    {
11        //Create ForkJoinPool using the default constructor.
12        ForkJoinPool pool = new ForkJoinPool();
13        //Create three FolderProcessor tasks. Initialize each one with a different
14        FolderProcessor system = new FolderProcessor("&quot;C:\\Windows&quot;;, &c
15        FolderProcessor apps = new FolderProcessor("&quot;C:\\Program Files&quot;;, &c
16        FolderProcessor documents = new FolderProcessor("&quot;C:\\Documents And
17        //Execute the three tasks in the pool using the execute() method.
18        pool.execute(system);
19        pool.execute(apps);
20        pool.execute(documents);
21        //Write to the console information about the status of the pool every se
22        //until the three tasks have finished their execution.
23        do
24        {
25            System.out.printf("&quot;*****\n&quot;");
26            System.out.printf("&quot;Main: Parallelism: %d\n&quot;;, pool.getParall
27            System.out.printf("&quot;Main: Active Threads: %d\n&quot;;, pool.getAct
28            System.out.printf("&quot;Main: Task Count: %d\n&quot;;, pool.getQueuedT
29            System.out.printf("&quot;Main: Steal Count: %d\n&quot;;, pool.getStealC
30            System.out.printf("&quot;*****\n&quot;");
31            try
32            {
33                TimeUnit.SECONDS.sleep(1);
34            } catch (InterruptedException e)
35            {
36                e.printStackTrace();
37            }

```



```

38     } while ((!system.isDone()) || (!apps.isDone()) || (!documents.isDone()))
39     //Shut down ForkJoinPool using the shutdown() method.
40     pool.shutdown();
41     //Write the number of results generated by each task to the console.
42     List<String> results;
43     results = system.join();
44     System.out.printf("System: %d files found.\n", results.size());
45     results = apps.join();
46     System.out.printf("Apps: %d files found.\n", results.size());
47     results = documents.join();
48     System.out.printf("Documents: %d files found.\n", results.size());
49 }
50 }

```

Output of above program will look like this:

Main: Parallelism: 2

Main: Active Threads: 3

Main: Task Count: 1403

Main: Steal Count: 5551

Main: Parallelism: 2

Main: Active Threads: 3

Main: Task Count: 586

Main: Steal Count: 5551

System: 337 files found.

Apps: 10 files found.

Documents: 0 files found.

How it works?

In the `FolderProcessor` class, Each task processes the content of a folder. As you know, this content has the following two kinds of elements:

- Files
- Other folders

If the task finds a folder, it creates another Task object to process that folder and sends it to the pool using the **fork()** method. This method sends the task to the pool that will execute it if it has a free worker-thread or it can create a new one. The **method returns immediately, so the task can continue processing** the content of the folder. For every file, a task compares its extension with the one it's looking for and, if they are equal, adds the name of the file to the list of results.

Once the task has processed all the content of the assigned folder, it waits for the finalization of all the tasks it sent to the pool using the **join()** method. This method called in a task waits for the finalization of its execution and returns the value returned by the **compute()** method. The task groups the results of all the tasks it sent with its own results and returns that list as a return value of the **compute()** method.

Difference between Fork/Join Framework And ExecutorService

The **main difference between the Fork/Join and the Executor frameworks is the work-stealing algorithm**. Unlike the Executor framework, when a task is waiting for the finalization of the sub-tasks it has created using the join operation, the thread that is executing that task (called worker thread) looks for other tasks that have not been executed yet and begins its execution. By this way, the threads take full advantage of their running time, thereby improving the performance of the application.

Existing Implementations in JDK

There are some generally useful features in Java SE which are already implemented using the fork/join framework.

1) One such implementation, introduced in Java SE 8, is used by the **java.util.Arrays class for its parallelSort() methods**. These methods are similar to **sort()**, but leverage concurrency via

the fork/join framework. Parallel sorting of large arrays is faster than sequential sorting when run on multiprocessor systems.

2) Parallelism used in **Stream.parallel()**. Read more about this [parallel stream operation in java 8](#).

Conclusion

Designing good multi-threaded algorithms is hard, and **fork/join doesn't work in every circumstance**. It's very useful within its own domain of applicability, but in the end, you have to decide whether your problem fits within the framework, and if not, **you must be prepared to develop your own solution**, building on the superb tools provided by java.util.concurrent package.

References

<http://gee.cs.oswego.edu/dl/papers/fj.pdf>

<http://docs.oracle.com/javase/tutorial/essential/concurrency/forkjoin.html>

<http://www.packtpub.com/java-7-concurrency-cookbook/book>

Happy Learning !!



Related Posts:

1. [Java executor framework tutorial and best practices](#)
2. [Difference between yield\(\) and join\(\) in threads in java?](#)
3. [Java 8: String join \(CSV\) example](#)
4. [Java IO : FileFilter Example Tutorial](#)
5. [Java 7 Changes, Features and Enhancements](#)
6. [When to use CountdownLatch : Java concurrency example tutorial](#)

■ FORK/JOIN FRAMEWORK

■ JAVA 7

PREVIOUS POST

[Difference between java.exe and javaw.exe](#)

NEXT POST

[Java 7 Changes, Features and Enhancements](#)

4 THOUGHTS ON “FORK/JOIN FRAMEWORK TUTORIAL: FORKJOINPOOL EXAMPLE”

emil

SEPTEMBER 27, 2014 AT 4:05 PM

Thank you so much. Simple and Nice explanation!

↩ REPLY

HIMANSU NAYAK

JUNE 5, 2014 AT 5:57 AM

Hi Lokesh,

“Where numberOfProcessors = Runtime.getRuntime().availableProcessors();”

no of processors can be either multiple core inside a single chip processor or single core multiple chip processor

↩ REPLY

★ Lokesh

JUNE 5, 2014 AT 6:54 AM

Any of both is possible. In my machine, it is first case.

↩ REPLY

Muhammed Shakir

SEPTEMBER 24, 2014 AT 4:25 AM

If hyper threading is enabled, then 4 cores will be capable of executing 8 threads at a time. So `Runtime.getRuntime().availableProcessors` will return 8. So number of processors is the ability of your processor architecture to execute the “number of threads at a time”.

↩ REPLY

Amazon Kinesis
Process real-time data
at massive scale.



Try Now



Note:- In comment box, please put your code inside `[java] ... [/java]` OR `[xml] ... [/xml]` tags otherwise it may not appear as intended.

LEAVE A REPLY

Your email address will not be published. Required fields are marked *

Name *

Email *

Website

Are you Human and NOT robot? Solve this to prove !! *

three + = 11

Comment

You may use these HTML tags and attributes:

 <abbr title=""> <acronym title=""> <blockquote cite=""> <cite> <code> <del datetime=""> <i> <q cite=""> <strike>

POST COMMENT

☐ Notify me of follow-up comments by email.

☒ Notify me of new posts by email.

SUBSCRIBE TO BLOG VIA EMAIL

Join 3,668 other subscribers

Email Address

SUBSCRIBE

REFERENCES

- Oracle Java Docs
- Spring 3 Reference
- Spring 4 References
- Jboss RESTEasy JAX-RS
- Hibernate Developer Guide
- JUnit Wiki
- Maven FAQs
- Dzone Dev Links
- JSP Homepage
- ANT Homepage

META LINKS

[Advertise](#)

[Share your knowledge](#)

[Privacy policy](#)

[Contact Us](#)

[About Us](#)

POPULAR POSTS

[Spring 3 and hibernate integration tutorial with example](#)

[Reading/writing excel files in java : POI tutorial](#)

[How hashmap works in java](#)

[Singleton design pattern in java](#)

[Useful java collection interview questions](#)

[Understanding hibernate first level cache with example](#)

[How hibernate second level cache works?](#)

[Working with hashCode and equals methods in java](#)

[How to make a java class immutable](#)

[4 ways to schedule tasks in Spring 3 : @Scheduled example](#)

© 2012-2015 Lokesh Gupta All Rights Reserved